

Computer Vision 2025

Project 10

Transformer-based Satellite Image and Segmentation Generation for Ground-to-Aerial Image Matching

Michele Capponi

Summary

- Introduction
- Dataset
- Proposed methods
- Image Generation (1-4)
- Segmentation (1-2)
- Image Matching
- Experimental results
- Future developments



Ground-to-aerial image matching

Matching the right satellite image for each streetview images is a very complex task due to **semantic** differences and **domain** adaptation

Streetview



Corresponding satellite



This project aims to **tackle this gap** without the need of more informations (like GPS) but only using the streetview image.

Dataset: CVUSA

The [version](#) I used of **Cross-view USA** (CVUSA) contains:

- **Streetview** images
- **Satellite** images
- Satellite **segmentation**
- Polarmap (segmented and not)

The original dataset can be found [here](#), a large dataset containing millions of pairs of ground-level and aerial/satellite images from across the United States.

Proposed architecture

Input Streetview



Satellite image generation

Transformed

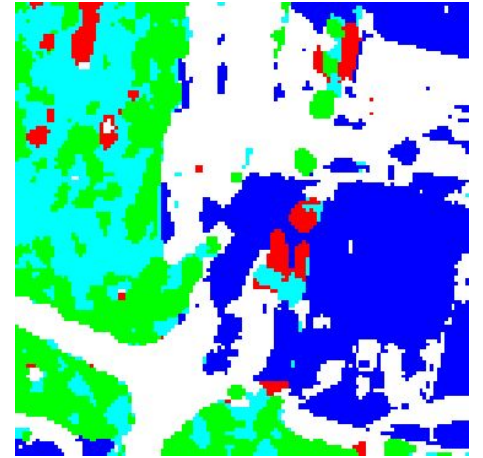


Joint Feature Learning Net

Feature Fusion Net

Matching

Segmentation



Candidate Satellite

Image Generation (1)

The main thing I tried were **Diffusion Transformers** I based my work on this [paper](#) by Meta

Scalable Diffusion Models with Transformers

William Peebles*
UC Berkeley

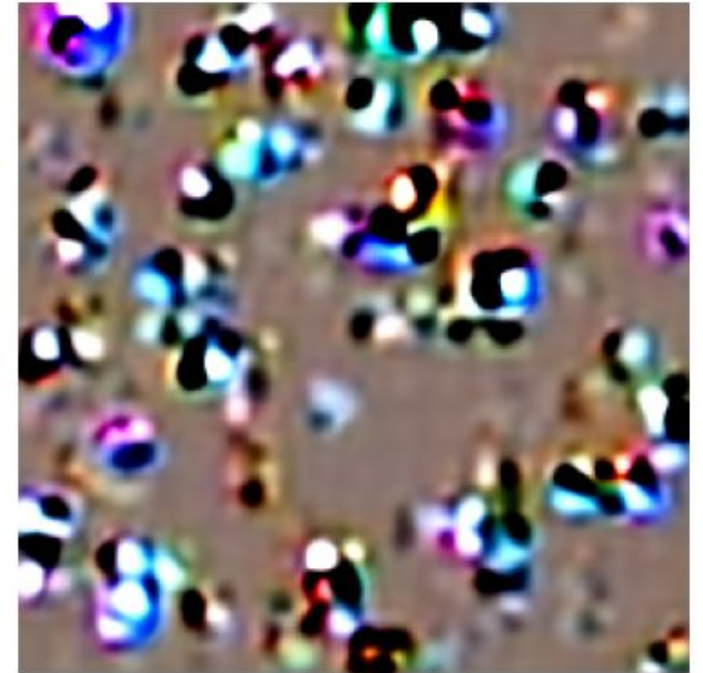
Saining Xie
New York University



DiT-XL/2 model, one of the biggest, with **VisualCLIP** and **stability-ai-sd-mse** VAE

Results?

Very **difficult conditioning**, as you can see...



Probably my main problem was the VAE since it **wasn't able to decode well its own encodings**

Image Generation (2)

I tried using a **VQGAN** encoder ([Taming-Transformers](#)), but the **checkpoints are not available** and I wasn't able to train my own.

Then I tried a **Unet** approach still using Diffusion. I opted for the **Stable Diffusion v1-v4** implementation (pre-trained on LAION-5B).

Reference: [huggingface](#)

- VisualCLIP encoding **Openai/clip-vit-base-patch32** (pre-trained on public images) → 10 epochs for last 2 layers
- **VAE default** of Stable Diffusion → Frozen
- **2 linear layer MLP projection** from CLIP to VAE

Image generation (3)

Pre-trained



Final Result



Original



iperparametri, loss
function, cazzate
varie e forse qualche
metrica? bo

- **AdamW**
- **ReduceLrOnPlateau**
- **Loss?**

Image generation (4)

Losses:

1. Noise loss: **MSE** ($1e-5$)
2. Clip: **1-CosineSimilarity** ($1e-6$)
3. VAE: **lpips**

In my last epoch this were my results:

Noise: 0.1791

Clip: 0.3957

VAE: 0.6177

Total: 1.5882

This means that the hardest task still is conditioning...

- in fact every time I regenerated an image I got a (substantially) **different output**
- good denoising but bad conditioning
- only learning truly **major features** (e.g. color, big roads or houses)


Segmentation (1)

I used **SegFormer** because there are available checkpoints pre-trained on [Ade20K](#) which has many spatial informations and domain closer to urban and aerial image.

The chosen model was [SegFormerB3](#).

It makes a significant **improvement** from B2 model **without having too many parameters** and it's quite close to B4 and B5 with much less parameters.

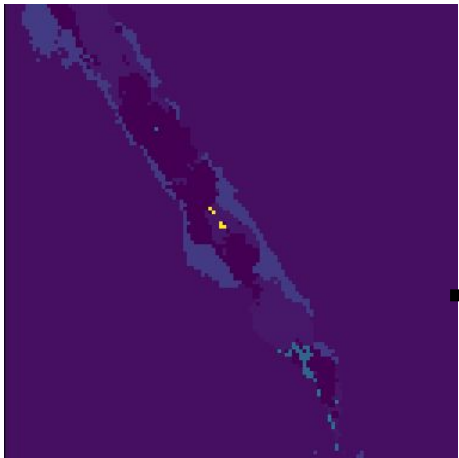
[Reference](#)



Segmentation (2)

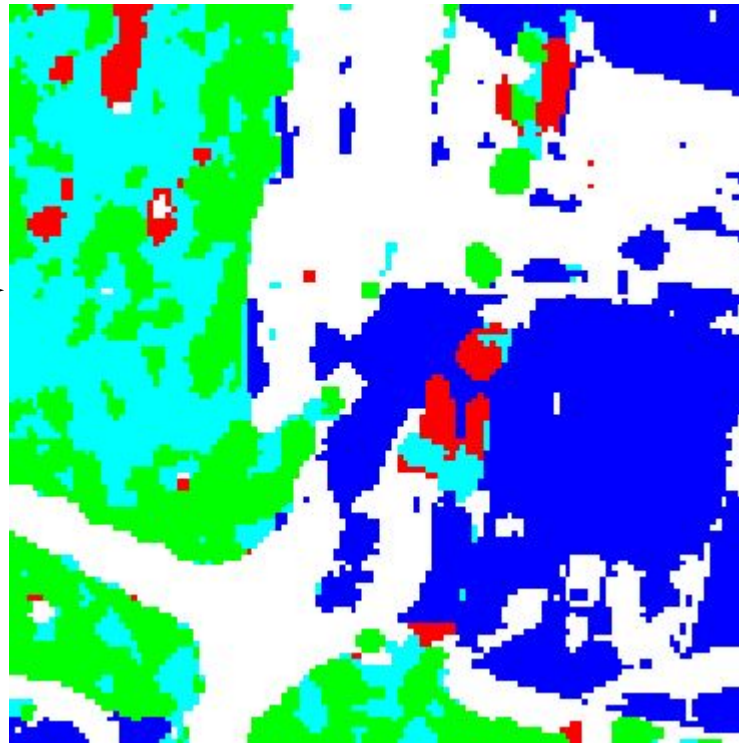
- `CrossEntropyLoss`
- `AdamW`
- `ReduceLrOnPlateau` ($\text{lr}=3\text{e-}5$)

Pre-trained

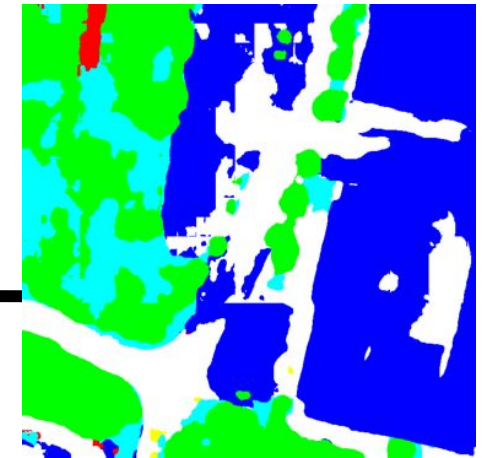


SegFormerB3 on
Ade20K:
mIoU = 0.49

Fine-tuned



Original



Fine-tuned on
CVUSA:
mIoU = 0.39

Feature extraction

JointFeatureLearningNet:

- 4 pre-trained VGG16 (Ground, Synthetic, Seg Synthetic, Candidate)
→ Joint representation

FeatureFusionNet:

- 1 FFN trained from scratch
→ Concatenate
- **Triplet Loss**
- **AdamW** (different lr for VGG and FFN)

Last epoch's loss: **0.15**

Metrics

Base accuracy: 51%

With 10 competitors:

Top1: 9%

Top5: 49%

For comparison (with more competitors)

[ALCOR LAB's SAN](#) (with 360 FOV)

Top1: 77%

Top5: 92%

[Regmi and Shah FFN](#)

Top1: 49%

It clearly is unfair to compare my task with those results but that's the **baseline** at the moment. My task of course was broader but nonetheless these are the results.

Conclusion

What's the take away from this experiment?

- Synthetic Satellite generation is a complex task, especially with bigger architectures like transformers
- The technique works but needs **more experimenting**
- It was **interesting and fun** to experiment with it
- Will be much **easier and more effective** in the future

Future Developments

Having more hardware capabilities (and time):

- Try again with **Diffusion Transformers** either fine-tuning the **VAE** or training it from scratch
- More fine-tuning on **VisualClip**
- More training on **SegFormer3**
- Adding **segmap** to the JointFeatureLearningNet
- Try with **polarmaps**
- **More complex training loop** for FFN and VGG
- Would be nice to see it reversed (**from satellite to streetview**, maybe embedding it in a satellite?)

Thanks for your attention

